

# In Defense of One-Vs-All Classification

*Journal of Machine Learning Research (2004)*

*Ryan Rifkin Aldebaro Klautau*

**Devin Conathan**  
**Josh Vahala**




**ECE 901: Large Scale Machine Learning  
and Optimization**



# Context: Multiclass Classification

Many machine learning classification problems are ***multiclass*** in nature

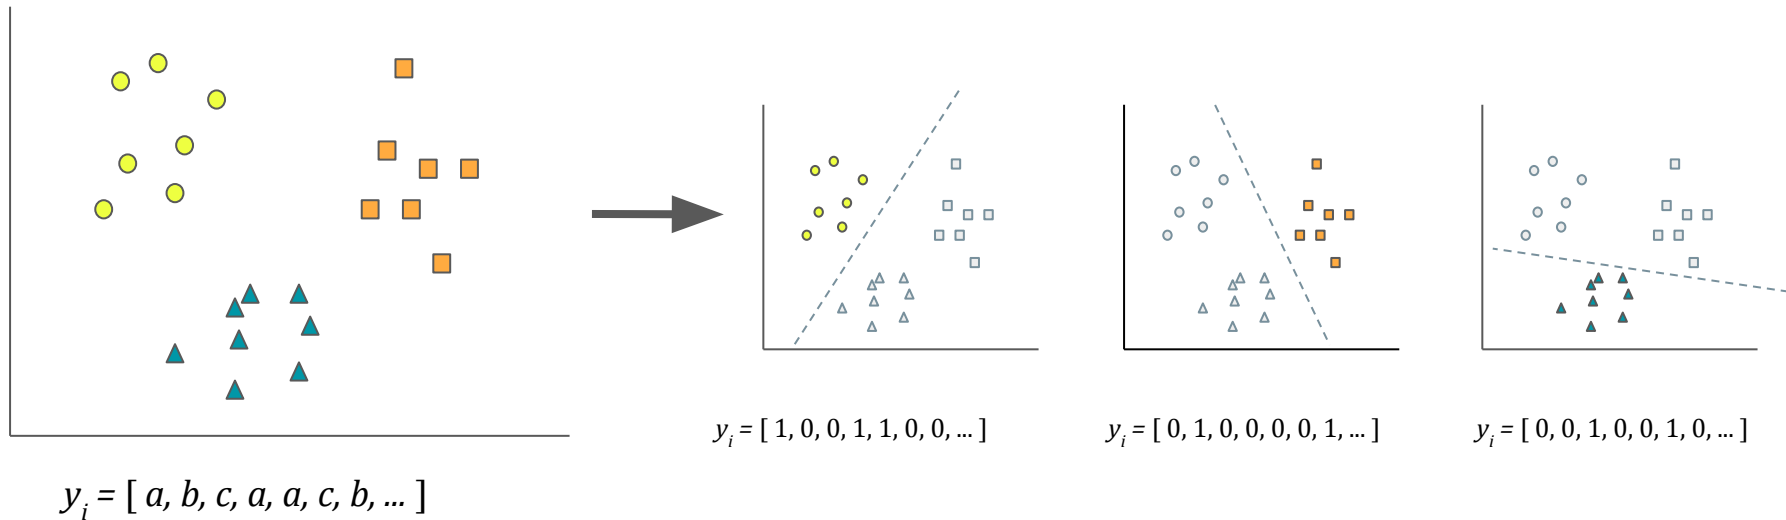
The loss function and optimization still fit our optimization framework



$$\min_f \sum_{i=1}^n L(f(x_i), y_i) + R(f)$$

# One vs. All Classification (OVA)

- About the simplest thing you could do
- For a multiclass problem with  $C$  classes: train  $C$  classifiers, one for each class
  - A new sample gets assigned class with the greatest predicted probability
- Simple to implement and embarrassingly parallelizable
- Python's scikit-learn provides a `OneVsRestClassifier` wrapper
- Alternatively, consider All vs. All Classification (AVA), where you train  $C(C-1)/2$  classifiers: one for each possible pair of classes



# But...

Some people decided that it's more fun to do things in way more complicated ways

This paper is essentially a literature review for these more complicated ways

The general theme is that the more complicated ways are:

- harder to implement
- slower to train and often computationally infeasible
- provide negligible (if any) performance boost

# Asymptotic Single-Machine Approach

## Theoretical motivation:

Derive a multiclass SVM that asymptotically behaves like the Bayes-optimal solution

If  $p(\mathbf{x})$  is the probability that  $\mathbf{x}$  is in some class, then:


the minimizer of  $E[(1 - yf(\mathbf{x}))_+]$  is  $f(\mathbf{x}) = \text{sign}(p(\mathbf{x}) - \frac{1}{2})$   
(the Bayes-optimal solution)

**One-VS-All SVM does NOT have this property:**

$$f_i(\mathbf{x}) \rightarrow \text{sign}(p_i(\mathbf{x}) - \frac{1}{2}) \text{ as } \ell \rightarrow \infty$$

$$\arg \max_i p_i(\mathbf{x}) \geq \frac{1}{2} \quad f_i(\mathbf{x}) = 1, \text{ and } f_j(\mathbf{x}) = -1 \text{ for } j \neq i \quad \checkmark$$

$$\arg \max_i p_i(\mathbf{x}) < \frac{1}{2} \quad f_i(\mathbf{x}) = -1 \quad \forall i \quad \times$$



# Method

Define a target vector  $v_i$  for  $1 \leq i \leq N$

$$v_i = \begin{bmatrix} -\frac{1}{N-1} \\ \vdots \\ 1 \\ \vdots \\ -\frac{1}{N-1} \end{bmatrix} \leftarrow i\text{th coordinate}$$

$$\begin{aligned} \min_{f_1, \dots, f_N \in \mathcal{H}_K} \quad & \frac{1}{\ell} \sum_{i=1}^{\ell} \sum_{j=1, j \neq y_i}^N (f_j(\mathbf{x}_i) + \frac{1}{N-1})_+ + \lambda \sum_{j=1}^C \|f_j\|_K^2 \\ \text{subject to :} \quad & \sum_{j=1}^C f_j(\mathbf{x}) = 0, \quad \forall \mathbf{x}. \end{aligned}$$

Then,  $f_i(\mathbf{x}) = 1$  if class  $i$  is the most likely and  $f_i(\mathbf{x}) = -\frac{1}{N-1}$  otherwise.

# Problems

- 1) Entirely asymptotic
  - a) Equivalent to other asymptotically accurate density estimation methods
  - b) With limited data, discriminative methods like SVM perform better
  - c) Ignores regularization
- 2) Overlapping class densities create additional issues
  - a) Classification accuracy is dependent on the most likely class in the region
  - b) High dimensional data will require many points to differentiate classes



# Multiclass Generalization of SVMs

Standard approach:                      find      $f(x) = \sum_{j=1}^{\ell} c_j K(\mathbf{x}, \mathbf{x}_j) + b.$

Proposed approach:                      find      $f_i(x) = \sum_{j=1}^{\ell} c_{ij} K(\mathbf{x}, \mathbf{x}_j) + b_i. \quad \leftarrow \text{find } N \text{ functions } \{f_1, f_2, \dots, f_N\} \text{ at once}$

***Instead of incurring cost for each machine (f), incur cost relative to values from other machines***


Vapnik and Blantz (1998), Weston and Watkins (1998)	Crammer and Singer (2001)
$\min_{\mathbf{f}_1, \dots, \mathbf{f}_N \in \mathcal{H}, \xi \in \mathbb{R}^{\ell(N-1)}} \quad \sum_{i=1}^N \ f_i\ _K^2 + C \sum_{i=1}^{\ell} \sum_{j \neq y_i} \xi_{ij}$ <p style="text-align: center;">subject to :    <math>f_{y_i}(\mathbf{x}_i) + b_{y_i} \geq f_j(\mathbf{x}_i) + b_j + 2 - \xi_{ij},</math>  <math>\xi_{ij} \geq 0.</math></p> <p><math>(N-1)\ell</math>     slack variables</p>	$\min_{\mathbf{f}_1, \dots, \mathbf{f}_N \in \mathcal{H}, \xi \in \mathbb{R}^{\ell}} \quad \sum_{i=1}^N \ f_i\ _K^2 + C \sum_{i=1}^{\ell} \xi_i$ <p style="text-align: center;">subject to :    <math>f_{y_i}(\mathbf{x}_i) \geq f_j(\mathbf{x}_i) + 1 - \xi_i,</math>  <math>\xi_i \geq 0.</math></p> <p><math>\ell</math>     slack variables</p>

where  $K(\mathbf{x}_1, \mathbf{x}_2) = \exp^{-\gamma \|\mathbf{x}_1 - \mathbf{x}_2\|^2},$  and  $\|f_i\|_K^2 = \mathbf{c}_i.^T K \mathbf{c}_i.,$



# Multiclass Generalization of SVMs

- 1) Claimed performance boost over OVA
  - a) Likely didn't tune OVA parameters well
  - b) Similar performance overall, but harder optimization problem



# Error-Correcting Code Approaches

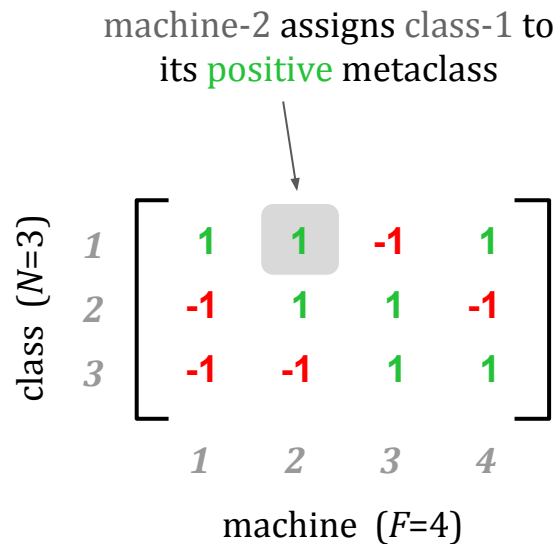
Define a matrix  $M \in \{-1, 1\}^{N \times F}$

where  $N$  is the number of classes and  $F$  is the number of machines

$j$ th machine solves 
$$\min \sum_{i=1}^{\ell} V(f_j(\mathbf{x}_i), M_{y_{ij}}) + \lambda \|f_j\|_K^2$$

To classify a new  $\mathbf{x}$ , calculate  $[f_1, f_2, \dots, f_F]$  and choose the class that minimizes the **Hamming distance** from the corresponding row of  $M$ :

$$f(\mathbf{x}) = \arg \min_{r \in 1, \dots, N} \sum_{i=1}^F \left( \frac{1 - \text{sign}(M_{ri} f_i(\mathbf{x}))}{2} \right)$$



# Improved Error-Correcting Code Approach

Expand the previous concept to allow zeros in  $M$ :

$$M \in \{-1, 0, 1\}^{N \times F}$$

This framework now encapsulates *OVA*, *AVA*, and general *Error-Correcting Code* classifiers

Also, introduce a more general **loss-based decoding** scheme, which significantly improves performance over Hamming distance:

$$f(\mathbf{x}) = \arg \min_{r \in 1, \dots, N} \sum_{i=1}^F L(M_{ri} f_i(\mathbf{x})).$$

$$\text{class (N=4)} \begin{bmatrix} 1 & -1 & -1 & -1 \\ -1 & 1 & -1 & -1 \\ -1 & -1 & 1 & -1 \\ -1 & -1 & -1 & 1 \end{bmatrix}$$

**One-vs-All**

$$\text{class (N=4)} \begin{bmatrix} 1 & 1 & 1 & 0 & 0 & 0 \\ -1 & 0 & 0 & 1 & 1 & 0 \\ 0 & -1 & 0 & -1 & 0 & 1 \\ 0 & 0 & -1 & 0 & -1 & -1 \end{bmatrix}$$

**All-vs-All**

# Error-Correcting Code Problems

1. Lots of tuning
  - a. For general error-correcting, not clear what matrix is the best
  - b. If you have many classes, large number of possibilities
2. Difference in performance is negligible
3. Results reported in *Allwein et al. (2000)* show better performance than OVA, but their OVA SVM classifiers were not tuned



# Theoretical Results

If the following conditions hold:

- Underlying classifier is a regularized least squares classifier (*RLSC*)
- The coding scheme's classifiers are independent
- The coding matrix contains no zeros

Then the problem reduces to an *OVA* multiclass classifier

i.e. the predictions will be ***exactly the same***

In particular, the *complete coding scheme*, which is the matrix that contains all unique  $\{+1, -1\}$  codes, has these properties.

# Experiments and Results

Name	OVA	AVA	COM	DEN	SPA
soybean-large	19	171	262143	43	64
letter	26	325	33554431	48	71
satimage	6	15	31	26	39
abalone	29	406	268435455	49	73
optdigits	10	45	511	34	50
glass	6	15	31	26	39
car	4	6	7	20	30
spectrometer	48	1128	1.407e+014	56	84
yeast	10	45	511	34	50
page-blocks	5	10	15	24	35

Table 5: Number of possible binary classifiers for each code matrix.

# Experiments and Results

Data Set	AVA	OVA	DIFF	AGREE	BOOTSTRAP
soybean-large	6.38	5.85	0.530	0.971	[-0.008, 0.019]
letter	3.85	2.75	1.09	0.978	[0.008, 0.015]
satimage	8.15	7.80	0.350	0.984	[-5E-4, 0.008]
abalone	72.32	79.69	-7.37	0.347	[-0.102, -0.047]
optdigits	3.78	2.73	1.05	0.982	[0.006, 0.016]
glass	30.37	30.84	-470	0.818	[-0.047, 0.037]
car	0.41	1.50	-1.09	0.987	[-0.016, -0.006]
spectrometer	42.75	53.67	-10.920	0.635	[-0.143, -0.075]
yeast	41.04	40.30	0.740	0.855	[-0.006, 0.021]
page-blocks	3.38	3.40	-0.020	0.991	[-0.002, 0.002]

Table 6: SVM test error rate (%), OVA vs. AVA.

Data Set	SPA	OVA	DIFF	AGREE	BOOTSTRAP
soybean-large	6.12	5.85	0.270	0.968	[-0.011, 0.016]
letter	3.55	2.75	0.800	0.980	[0.005, 0.011]
satimage	8.85	7.80	1.05	0.958	[0.003, 0.018]
abalone	75.67	79.69	-4.02	0.352	[-0.067, -0.014]
optdigits	3.01	2.73	0.280	0.984	[-0.002, 0.008]
glass	28.97	30.84	-1.87	0.738	[-0.070, 0.033]
car	0.81	1.50	-0.69	0.988	[-0.011, -0.003]
spectrometer	52.73	53.67	-0.940	0.744	[-0.038, 0.019]
yeast	40.16	40.30	-0.140	0.855	[-0.015, 0.013]
page-blocks	3.84	3.40	0.440	0.979	[0.001, 0.007]

Table 8: SVM test error rate (%), OVA vs. SPARSE.

Data Set	DEN	OVA	DIFF	AGREE	BOOTSTRAP
soybean-large	5.58	5.85	-0.270	0.963	[-0.019, 0.013]
letter	2.95	2.75	0.200	0.994	[5E-4, 0.004]
satimage	7.65	7.80	-0.150	0.985	[-0.006, 0.003]
abalone	73.18	79.69	-6.51	0.393	[-0.092, -0.039]
optdigits	2.61	2.73	-0.12	0.993	[-0.004, 0.002]
glass	29.44	30.84	-1.40	0.911	[-0.042, 0.014]
car	-	1.50	-	-	-
spectrometer	54.43	53.67	-0.760	0.866	[-0.011, 0.026]
yeast	40.30	40.30	0.00	0.900	[-0.011, 0.011]
page-blocks	-	3.40	-	-	-

Table 7: SVM test error rate (%), OVA vs. DENSE.

Data Set	COM	OVA	DIFF	AGREE	BOOTSTRAP
soybean-large	-	5.85	-	-	-
letter	-	2.75	-	-	-
satimage	7.80	7.80	0.00	0.999	[-1E-3, 1E-3]
abalone	-	79.69	-	-	-
optdigits	2.67	2.73	-0.060	0.996	[-0.003, 0.002]
glass	29.44	30.84	-1.340	0.911	[-0.042, 0.014]
car	1.68	1.50	-0.180	0.998	[5.79E-4, 0.003]
spectrometer	-	53.67	-	-	-
yeast	38.61	40.30	-1.690	0.906	[-0.028, -0.005]
page-blocks	3.49	3.40	-0.090	0.983	[-0.002, 0.004]

Table 9: SVM test error rate (%), OVA vs. COMPLETE.

# Takeaways

- 1) OVA and AVA are very simple to implement and perform well
- 2) OVA is not significantly outperformed by proposed multiclass methods
  - a) This does not mean there isn't a method that will perform better!
- 3) AVA has a speed advantage to OVA because fewer examples per optimization, but requires more trained classifiers



# Questions?